

# Structure of Deep Learning Frameworks: computational graph, autodiff, and optimizers

Joris Mollinga  
SURF




# Computational Graphs

- Represents math in graph format
- Nodes and edges


# Computational Graphs

- Represents math in graph format
- Nodes and edges




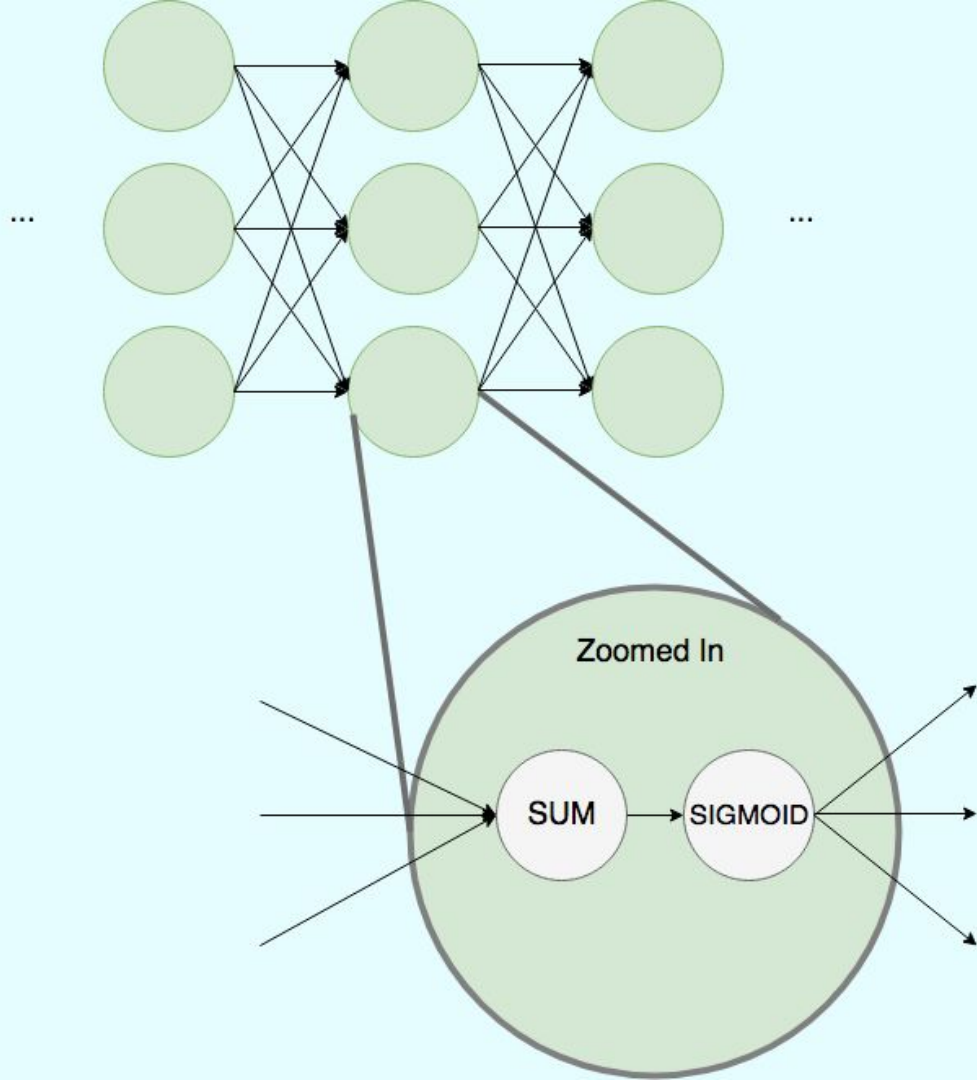
# Computational Graphs

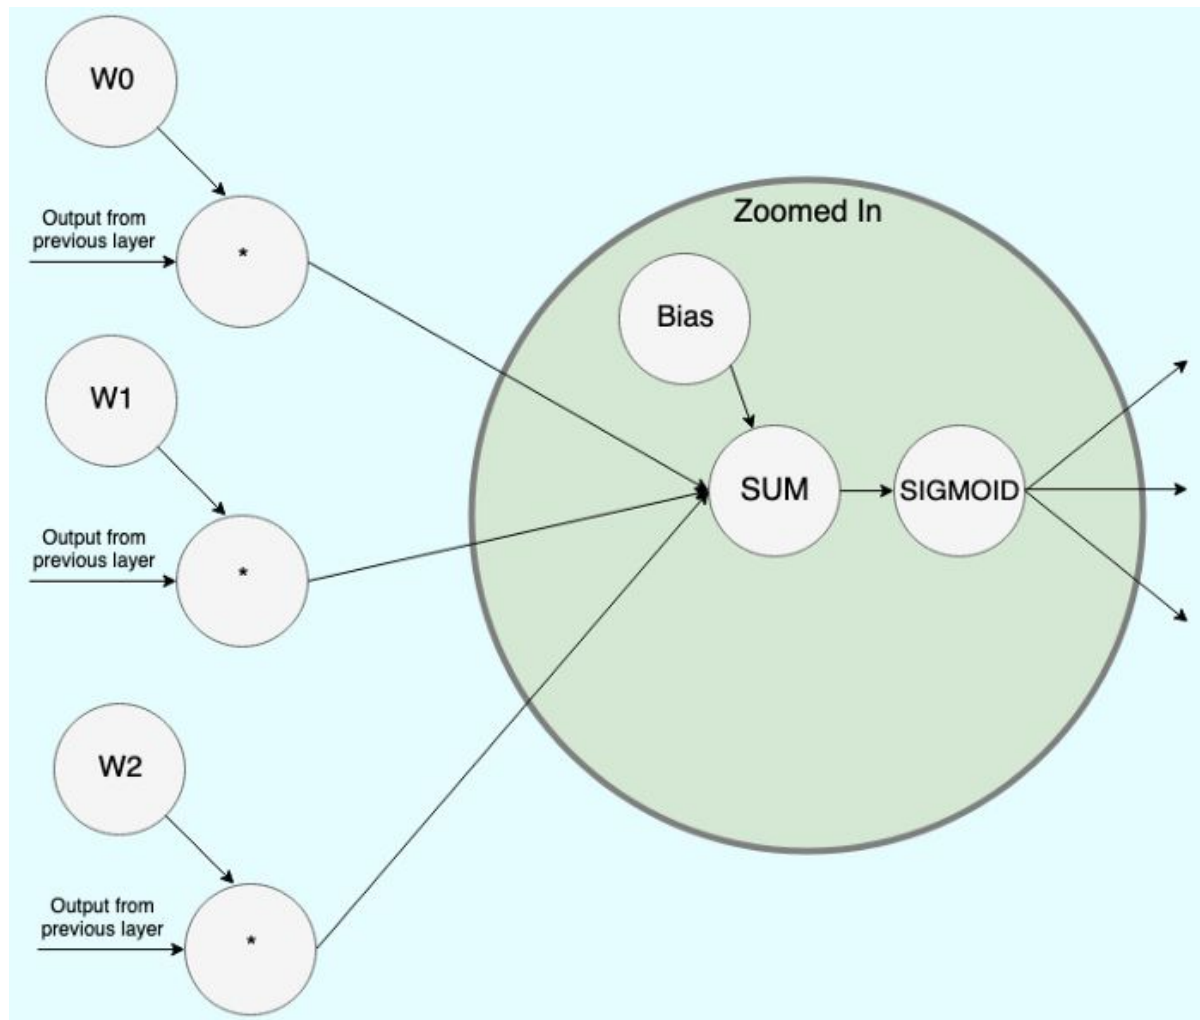
- Represents math in graph format
- Nodes and edges



$$f(x, y) = ax^2 + bxy + cy^2$$







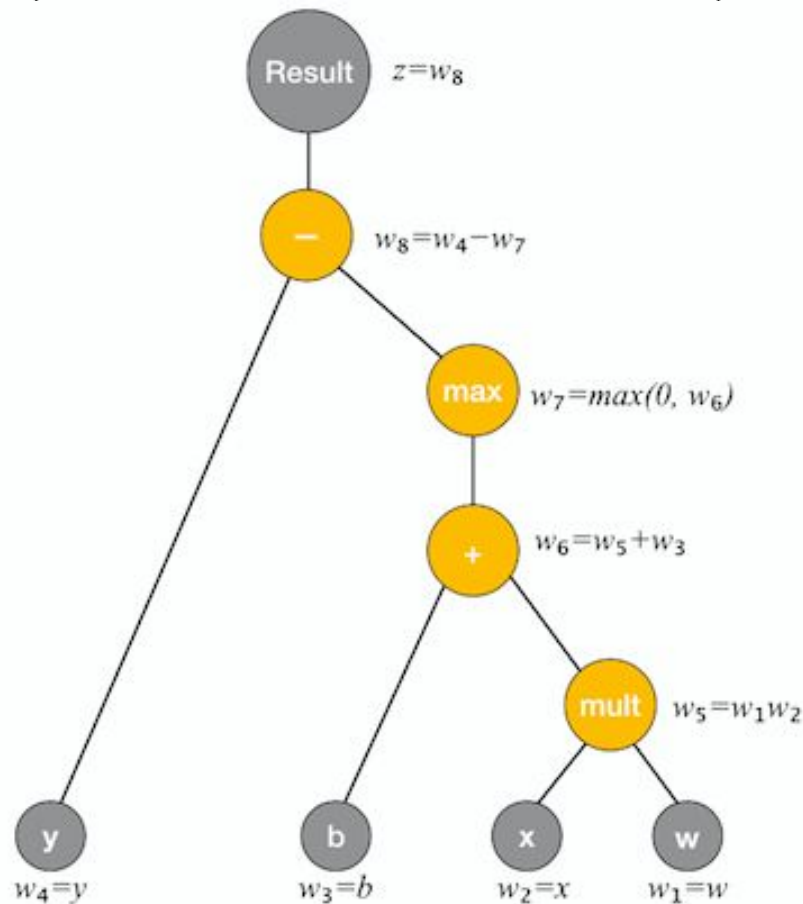


## Autodiff

$$C(y, w, x, b) = y - \max(0, w \cdot x + b)$$




$$C(y, w, x, b) = y - \max(0, w \cdot x + b)$$




Node	Expression	Value
$w_1$	$w$	2
$w_2$	$x$	1
$w_3$	$b$	1
$w_4$	$y$	5
$w_5$	$w_1 \cdot w_2$	2
$w_6$	$w_5 + w_3$	3
$w_7$	$\max(0, w_6)$	3
$w_8$	$w_4 - w_7$	2
$z$	$w_8$	2

$$C(y, w, x, b) = y - \max(0, w \cdot x + b)$$



$\frac{\delta w_5}{\delta w_1} = \frac{\delta w_1 w_2}{\delta w_1} = w_2$	$\frac{\delta w_5}{\delta w_2} = \frac{\delta w_1 w_2}{\delta w_2} = w_1$
$\frac{\delta w_6}{\delta w_5} = \frac{\delta w_5 + w_3}{\delta w_5} = 1$	$\frac{\delta w_6}{\delta w_3} = \frac{\delta w_5 + w_3}{\delta w_3} = 1$
$\frac{\delta w_7}{\delta w_6} = \frac{\delta \max(0, w_6)}{\delta w_6} = \begin{cases} 0, & x < 0 \\ 1, & x > 0 \end{cases}$	$\frac{\delta w_8}{\delta w_7} = \frac{\delta w_4 - w_7}{\delta w_7} = -1$
$\frac{\delta w_8}{\delta w_4} = \frac{\delta w_4 - w_7}{\delta w_4} = 1$	$\frac{\delta z}{\delta w_8} = 1$



$$\frac{\delta z}{\delta w_1} = \frac{\delta z}{\delta w_8} \times \frac{\delta w_8}{\delta w_7} \times \frac{\delta w_7}{\delta w_6} \times \frac{\delta w_6}{\delta w_5} \times \frac{\delta w_5}{\delta w_1} = 1 \times (-1) \times \begin{cases} 0, & x < 0 \\ 1, & x > 0 \end{cases} \times 1 \times w_2 = \begin{cases} 0, & x < 0 \\ 1, & x > 0 \end{cases}$$



# Frameworks

- Caffe
- NVCaffe
- IntelCaffe
- PyTorch
- PyTorch Lightning
- Tensorflow
- Horovod (DL distribution framework only)



# PyTorch

- Probably the most popular package nowadays
- Multi-GPU support (more than one node)
- Supports cuDNN and MKL-DNN
- Supports various communication backends: Gloo, MPI, NCCL



# Tensorflow

- Also very popular
- Multi-GPU in one node through *tf.device*
- Parallelism across nodes with *tf.distribute*
- Support for TPU
- Supports cuDNN and MKL-DNN
- Supports NCCL allreduce
- For CPU, build from source





# Tensorflow optimization tips

- Use *TfRecords* to prevent I/O bottlenecks
- Overlap computation and data preparation using *tf.data.Dataset.prefetch*
- Parallelize data transformation using *tf.data.Dataset.map*
- If data fits in memory, use *tf.data.Dataset.cache*



# Horovod

Is a distribution framework for deep learning (not a deep learning framework itself). Design goals:

- Minimal code changes to make serial program distributed
- High performance distribution



# Horovod

- Support for Keras, MXNet, Tensorflow and PyTorch
- Supports MPI and MCCL as communication backends
- Requires about 6 lines of code changes
- Has it's own profiling ability, making it easy to assess communication overhead.



# PyTorch Lightning

- PyTorch wrapper for high performance AI research
- Supports various distribution strategies (horovod, data parallel, model parallel)
- Has other convenient features too.

**Thank you!**

---